

Should you believe in the Shanghai ranking?

An MCDM view

Jean-Charles Billaut · Denis Bouyssou · Philippe Vincke

Received: 26 May 2009 / Published online: 2 December 2009
© Akadémiai Kiadó, Budapest, Hungary 2009

Abstract This paper proposes a critical analysis of the “Academic Ranking of World Universities”, published every year by the Institute of Higher Education of the Jiao Tong University in Shanghai and more commonly known as the *Shanghai ranking*. After having recalled how the ranking is built, we first discuss the relevance of the criteria and then analyze the proposed aggregation method. Our analysis uses tools and concepts from Multiple Criteria Decision Making (MCDM). Our main conclusions are that the criteria that are used are not relevant, that the aggregation methodology is plagued by a number of major problems and that the whole exercise suffers from an insufficient attention paid to fundamental structuring issues. Hence, our view is that the Shanghai ranking, in spite of the media coverage it receives, does not qualify as a useful and pertinent tool to discuss the “quality” of academic institutions, let alone to guide the choice of students and family or to promote reforms of higher education systems. We outline the type of work that should be undertaken to offer sound alternatives to the Shanghai ranking.

Keywords Shanghai ranking · Multiple criteria decision analysis · Evaluation models · Higher education

This paper is an abridged version of Billaut et al. (2009)

J.-C. Billaut
Laboratoire d’Informatique, Université François Rabelais Tours,
64 Avenue Jean Portalis, 37200 Tours, France
e-mail: jean-charles.billaut@univ-tours.fr

D. Bouyssou (✉)
CNRS–LAMSAD, UMR 7024, Université Paris Dauphine,
Place du Maréchal de Lattre de Tassigny, 75775 Paris Cedex 16, France
e-mail: bouyssou@lamsade.dauphine.fr

P. Vincke
Université Libre de Bruxelles, 50 avenue F. D. Roosevelt, CP. 130, 1050 Bruxelles, Belgium
e-mail: pvincke@ulb.ac.be

Introduction

In 2003, a group of people belonging to the *Institute of Higher Education* from the *Jiao Tong University* in Shanghai published on their web site their first “Academic Ranking of World Universities” (ARWU 2003–2009), more commonly known as the *Shanghai ranking*.¹ This ranking consisted in an ordered list of 500 universities in the whole world. Since then, this group publishes each year an updated version of the ranking. A description of the ranking can be found in Liu and Cheng (2005), whereas the story behind its creation is detailed in Liu (2009).

This ranking was almost immediately the subject of an extraordinary media coverage. Not only political decision makers used the results of the ranking so as to promote reforms of higher education systems but many academic institutions began to use their position in this ranking in their institutional communication. Apparently, this looks like a true success story.

Yet, almost immediately after the release of the first ranking, this enterprise has been the subject of fierce attacks. One of the earlier one was due to van Raan (2005a), which started a vigorous exchange with the authors of the Shanghai ranking (Liu et al. 2005; van Raan 2005b). Since then the attacks have been numerous and strong both in the academic literature (Buela-Casal et al. 2007; Dill and Soo 2005; Gingras 2008; Ioannidis et al. 2007; van Raan 2006; Vincke 2009; Zitt and Filliatreau 2006) and in various reports and position papers (Bourdin 2008; Brooks 2005; HEFCE 2008; Dalsheimer and Despréaux 2008; Desbois 2007; Kivinen and Hedman 2008; Kävelmark 2007; Marginson 2007; Saisana and D’Hombres 2008; Stella and Woodhouse 2006).² In views of such attacks, one could have expected a sharp decrease in the popularity of the Shanghai ranking. It could even have triggered its authors to stop publishing it. Quite the contrary happened. Each year, a new version of the ranking is released and each year the media coverage of the ranking seems to increase. Moreover, projects of transformation of higher education systems often appeal to the results of the Shanghai ranking. For instance, the present French Minister of Research and Higher Education was given by the French President the mission “to have two institutions among the world top 20 and ten among the world top 100”.³

This paper wishes to be a contribution to the analysis of the strengths and weaknesses of the Shanghai ranking. Our point of view will be that of persons specialized in Operations Research and having worked in the field of evaluation and decision models with multiple criteria (Bouyssou et al. 2000, 2006; T’kindt and Billaut 2006), while most of the previous analyses of the Shanghai ranking have concentrated on bibliometric aspects, starting with the important contribution of van Raan (2005b).

This paper is organized as follows. In “How the Shanghai ranking is built?”, we will briefly describe how the authors of the Shanghai ranking operate. In “An analysis of the criteria” will discuss the various criteria that are used. In “An MCDM view on the Shanghai ranking” will present a Multiple Criteria Decision Making (MCDM) view on the Shanghai ranking. A final section will discuss our findings.

¹ Since then, the authors of the Shanghai ranking have also produced, starting in 2007, a ranking of institutions distinguishing 5 different fields within Science, see <http://www.arwu.org/ARWU-FIELD2008.htm>. Since the methodology for these “field rankings” is quite similar to the one used for the “global ranking” analyzed in this paper, we will not further analyze them here.

² Furthermore, several special issues of the journal *Higher Education in Europe* have been devoted to the debate around university rankings

³ Letter dated 5 July 2007, our translation from French, source <http://www.elysee.fr/>, last accessed 18 September 2009. Unless otherwise stated, all URL mentioned below have been accessed at this date.

How the Shanghai ranking is built?

This section describes how the Shanghai ranking is built, based upon ARWU (2003–2009) and Liu and Cheng (2005). We concentrate on the last edition of the ranking published in 2008⁴ although the methodology has varied over time.

Who are the authors of the ranking?

The authors of the Shanghai ranking⁵ are a group of people belonging to the Institute of Higher Education of the Jiao Tong University in Shanghai. This group is headed by Professor Nian Cai Liu. According to Liu (2009), this group started its work on the ranking of universities in 1998, following an initial impulse of the Chinese government (the “985 Project” referred to in Liu 2009).

The authors of the ranking admit (Liu et al. 2005, p. 108) that they had, when they started to work on the subject, no particular expertise in bibliometry. They claim (Liu and Cheng 2005, p. 135) that they receive no particular funding for producing the ranking and that they are guided mainly by academic considerations. This is at variance with the situation for rankings such as the one produced by the Times Higher Education Supplement (2008).

The announced objective of the authors of the ranking is to have a tool allowing them understand the gap between Chinese universities and “world-class universities”, with the obvious and legitimate aim of reducing this gap. Because of the difficulty to obtain “internationally comparable data”, they decided to rank order universities based on “academic or research performance” (Liu and Cheng 2005, p. 133).

How were the universities selected?

The authors claim to have analyzed around 2000 institutions worldwide (Liu and Cheng 2005, p. 127–128). This is supposed to include all institutions having Nobel prize and Fields medal laureates, a significant number of papers published in *Nature* or *Science*, of highly cited researchers as given by Thomson Scientific (formerly ISI), and a significant amount of publications indexed in the Thomson Scientific databases. The authors of the ranking claim that this includes all major universities in each country. The published ranking only includes 500 institutions. The first 100 are ranked ordered. The remaining ones are rank by groups of 50 (till the 201th position) and then 100.

The criteria

The authors use six criteria belonging to four distinct domains. They are presented below.

Quality of education

This domain uses a single criterion: the number of alumni of the institution having received a Nobel prize (Peace and Literature are excluded, the Bank of Sweden prize in Economics included) or a Fields medal. An alumni is defined as a person having obtained a Bachelor, a

⁴ See [http://www.arwu.org/rank2008/ARWU2008Methodology\(EN\).htm](http://www.arwu.org/rank2008/ARWU2008Methodology(EN).htm). The 2009 edition of the ranking is scheduled to be released in November 2009.

⁵ We will often simply refer to them in this paper as “the authors of the ranking”.

Master or a Doctorate in the institution. If a laureate has obtained a degree from several institutions, each one receives a share. All prizes and medals do not have the same weight: they are “discounted” using a simple linear scheme (an award received after 1991 counts for 100%, an award received between 1981 and 1990 counts for 90%,...). When several persons are awarded the prize or the medal, each institution receives a share. This defines the first criterion labeled ALU.

Quality of faculty

This domain has two criteria. The first one counts the number of academic staff from the institution having received a Nobel prize (with the same definition as above) or a Fields medal. The conventions for declaring that someone is a member of the “academic staff” of an institution remain fuzzy. The following discounting scheme is applied: 100% for winners in after 2001, 90% for winners in 1991–2000, 80% for winners in 1981–1990,..., 10% for winners in 1911–1920. The case of multiple winners is treated as with criterion ALU. When a person has several affiliations, each institution receives a share. This defines criterion AWA.

The second criterion in this domain is the number of highly cited researchers in each of the 21 areas of Science identified by Thomson Scientific. These highly cited researchers, in each of the 21 domains, consist in a list of 250 persons who have received the largest number of citations in the domain according to the Thomson Scientific databases (see <http://www.isihighlycited.com>). This is computed over a period of 20 years. This defines criterion HiCi.

Research output

This domain has two criteria. The first one is the number of papers published in *Nature* and *Science* by the members of the academic staff of an institution during the last 5 years. This raises the problem of processing papers having multiple authors. The rule here is to give a weight of 100% to the corresponding author affiliation, 50% for first author affiliation (second author affiliation if the first author affiliation is the same as corresponding author affiliation), 25% for the next author affiliation, and 10% for other author affiliations. This defines criterion N&S. Since this criterion is little relevant for institutions specialized in Social and Human Sciences, it is “neutralized” for them.

The second criterion counts the number of papers published by the members of the academic staff of an institution. This count is performed using Thomson Scientific databases over a period of one year. Since it is well known that the coverage of the Thomson Scientific databases is not satisfactory for Social and Human Sciences, a coefficient of 2 is allocated to each publication indexed in the Social Science Citation Index. This defines criterion PUB.

Productivity

This domain has a single criterion. It consists in the “total score of the above five indicators divided by the number of Full Time Equivalent (FTE) academic staff” (Liu and Cheng 2005, p. 129). This criterion is “ignored” when the number of FTE academic staff could not be obtained.⁶ This defines criterion Py.

⁶ In ARWU (2003–2009), the authors of the ranking say that this number was obtained for “institutions in USA, UK, France, Japan, Italy, China, Australia, Netherlands, Sweden, Switzerland, Belgium, South Korea,

Data collection

Data are mostly collected on the web. This involves the official site of the Nobel Prizes (http://www.nobelprize.org/nobel_prizes/), the official site of the International Mathematical Union (<http://www.mathunion.org/general/prizes>, and various Thomson Scientific sites (<http://www.isihighlycited.com> and <http://www.isiknowledge.com>). The authors of the ranking (ARWU 2003–2009) do not exactly specify the source of the data for the number of FTE academic staff of each institution.⁷ The data used by the authors of the ranking are not made publicly available.

Normalization and aggregation

Each of the above six criteria is measured by a positive number. Each criterion is then normalized as follows. A score of 100 is given to the best scoring institution and all other scores are normalized accordingly. This leads to a score between 0 and 100 for each institution.

The authors say that “adjustments are made” when the statistical analyses reveal “distorting effects” (Liu and Cheng 2005, p. 129). The nature and the scope of these adjustments are not made public (Florian 2007, shows that these adjustments are nevertheless important).

The authors use a weighted sum to aggregate these normalized scores. The weights of the six criteria are ALU: 10%, AWA: 20%, N&S: 20%, HiCi: 20%, PUB: 20%, and Py: 10%. Hence each institution receives a score between 0 and 100. The final scores are then normalized again so that the best institution receives a score of 100. This final normalized score is used to rank order the institutions.

The 2008 results

Table 1 gives the list of the best 20 universities in the world according to the 2008 edition of the Shanghai ranking. Table 2 does the same for European universities.

A cursory look at Table 1 reveals the domination of US universities in the ranking. Within Europe, the domination of the UK is striking.

Figure 1 shows the distribution of the global normalized score for the 500 institutions included in the Shanghai ranking. The curve becomes very flat as soon as one leaves the top 100 institutions.

We would like to conclude this brief presentation of the Shanghai ranking by giving three quotes taken from Liu and Cheng (2005, p. 135). The Shanghai ranking uses “carefully selected objective criteria”, is “based on internationally comparable data that everyone can check”, and is such that “no subjective measures were taken”.

Footnote 6 continued

Czech, Slovenia, New Zealand, etc.”. We do not know if this means that this number was obtained for *all* institutions in these countries and only for them.

⁷ More precisely, they mention in ARWU (2003–2009) that this number was obtained “from national agencies such as National Ministry of Education, National Bureau of Statistics, National Association of Universities and Colleges, National Rector’s Conference”.

Table 1 The best 20 universities in the world in the Shanghai ranking (2008)

Rank	Institution	Country	ALU	AWA	HiCi	N&S	PUB	Py	Score
1	Harvard	USA	100.0	100.0	100.0	100.0	100.0	74.1	100.0
2	Stanford	USA	40.0	78.7	86.6	68.9	71.6	66.9	73.7
3	UC Berkeley	USA	69.0	77.1	68.8	70.6	70.0	53.0	71.4
4	Cambridge	UK	90.3	91.5	53.6	56.0	64.1	65.0	70.4
5	MIT	USA	71.0	80.6	65.6	68.7	61.6	53.9	69.6
6	CalTech	USA	52.8	69.1	57.4	66.1	49.7	100.0	65.4
7	Columbia	USA	72.4	65.7	56.5	52.3	70.5	46.6	62.5
8	Princeton	USA	59.3	80.4	61.9	40.5	44.8	59.3	58.9
9	Chicago	USA	67.4	81.9	50.5	39.5	51.9	41.3	57.1
10	Oxford	UK	59.0	57.9	48.4	52.0	66.0	45.7	56.8
11	Yale	USA	48.5	43.6	57.0	55.7	62.4	48.7	54.9
12	Cornell	USA	41.5	51.3	54.1	52.3	64.7	40.4	54.1
13	UC Los Angeles	USA	24.4	42.8	57.4	48.9	75.7	36.0	52.4
14	UC San Diego	USA	15.8	34.0	59.7	53.0	66.7	47.4	50.3
15	U Pennsylvania	USA	31.7	34.4	58.3	41.3	69.0	39.2	49.0
16	U Wash Seattle	USA	25.7	31.8	53.1	49.5	74.1	28.0	48.3
17	U Wisc Madison	USA	38.4	35.5	52.6	41.2	68.1	28.8	47.4
18	UC San Francisco	USA	0.0	36.8	54.1	51.5	60.8	47.5	46.6
19	Tokyo Univ	Japan	32.2	14.1	43.1	51.9	83.3	35.0	46.4
20	Johns Hopkins	USA	45.8	27.8	41.3	48.7	68.5	24.8	45.5

Source: ARWU (2003–2009)

An analysis of the criteria

We start our analysis of the Shanghai ranking by an examination of the six criteria that it uses. This analysis mainly wishes to be a synthesis of the already large literature on the subject that was mentioned in “[Introduction](#)” and makes much use of van Raan (2005a). However, whereas the reviewed literature mainly draws on bibliometric considerations, we will also be influenced by the literature on MCDM on the structuration of objectives and the construction of criteria.

Criteria linked to Nobel prizes and Fields medals

Two criteria (ALU and AWA) are linked to a counting of Nobel prizes and of Fields medals. These two criteria are especially problematic.

Let us first observe that, for criterion AWA, the prize and medals are attributed to the hosting institution *at the time of the announcement*. This may not be too much of a problem for Fields medals (they are only granted to people under 40). This is however a major problem for Nobel prizes. Indeed, a close examination of the list of these prizes reveals that the general rule is that the prize is awarded long after the research leading to it has been conducted. A classic example of such a situation is Albert Einstein. He conducted his research while he was employed by the Swiss Patent Office in Zurich. He received the Nobel Prize long after, while he was affiliated to the University of Berlin. Even when the winner of a Nobel prize has not moved, the lag between the time at which the research was

Table 2 The best 20 universities in Europe in the Shanghai ranking (2008)

Rank	Institution	Country	ALU	AWA	HiCi	N&S	PUB	Py	Score
4	Cambridge	UK	90.3	91.5	53.6	56.0	64.1	65.0	70.4
10	Oxford	UK	59.0	57.9	48.4	52.0	66.0	45.7	56.8
22	U Coll London	UK	31.2	32.2	38.6	44.3	65.8	35.4	44.0
24	Swiss Fed Inst Tech—Zurich	Switzerland	35.9	36.3	36.1	38.1	53.6	56.0	43.1
27	Imperial Coll	UK	18.6	37.4	39.9	38.2	61.8	39.4	42.4
40	U Manchester	UK	24.4	18.9	28.2	28.3	60.5	30.4	33.6
42	U Paris 06	France	36.6	23.6	23.1	27.3	58.2	21.3	33.1
45	U Copenhagen	Denmark	27.4	24.2	26.3	25.4	54.5	33.4	33.0
47	U Utrecht	Netherlands	27.4	20.9	28.2	28.8	53.3	26.0	32.4
49	U Paris 11	France	33.3	46.2	14.6	20.4	47.0	23.1	32.1
51	Karolinska Inst Stockholm	Sweden	27.4	27.3	31.8	18.3	50.1	25.7	31.6
53	U Zurich	Switzerland	11.2	26.8	24.7	27.5	50.2	32.4	31.0
55	U Edinburgh	UK	20.2	16.7	26.3	32.3	49.7	30.0	30.8
55	U Munich	Germany	33.1	22.9	16.3	25.6	52.7	31.8	30.8
57	TU Munich	Germany	41.1	23.6	25.3	18.9	44.8	30.6	30.5
61	U Bristol	UK	9.7	17.9	28.2	28.1	47.8	33.5	29.5
64	U Oslo	Norway	23.1	33.4	17.9	17.0	46.7	29.8	29.0
67	U Heidelberg	Germany	17.7	27.2	17.9	20.4	49.2	29.3	28.4
68	U Helsinki	Finland	16.8	17.9	21.9	20.8	53.8	30.1	28.3
70	Moscow State U	Russia	49.1	34.2	0.0	8.3	53.2	33.4	28.1

Source: ARWU (2003–2009)




Fig. 1 Distribution of normalized scores for the 500 universities in the Shanghai ranking (2008). Source: ARWU (2003–2009)

conducted and the time at which the award was announced is such that the criterion captures more the past qualities of an institution than its present research potential. Therefore, it does not seem unfair to say that the link between AWA and the quality of research conducted in an institution is, at best, extremely approximative. The same is clearly true for criterion ALU since the time lag is here even longer. At best, this criterion might give some indication on the ability an institution had several decades ago to give extremely bright people a stimulating environment. It has little to do with the present ability of an institution to provide an excellent education to its students.

One may also wonder why prizes attributed long ago are linked with the present quality of an institution. Although the discounting scheme adopted by the authors of the ranking tends to limit the impact of these very old prizes and medals, they still have some effect. Moreover, the discounting scheme that is adopted is completely arbitrary (e.g., why use a linear and not an exponential scheme?)

The options taken by the authors of the ranking on these two criteria involve many other biases. A bias in favor of countries having known few radical political changes since 1901. A bias towards institutions having been created long ago and having kept the same name throughout their history. This is not the case for most institutions in continental Europe (think of the many wars and radical political changes that have happened in Europe since 1901). This has led to really absurd situations. For instance the two universities (Free university of Berlin and Humboldt University, using their names in English) created in Berlin after the partition of Germany and, therefore, the splitting of the University of Berlin, quarrelled over which one should get the Nobel Prize of Albert Einstein (see Enserink 2007, on this astonishing case). It turned out that depending on the arbitrary choice of the university getting this prize, these two institutions had markedly different positions in the ranking. Unfortunately, Germany is not an isolated example as revealed by our analysis of the French case in Billaut et al. (2009).

Finally, as pointed to us by a referee, these two criteria are based on Prizes and Medals that are far from covering all important scientific fields. Distinctions such as the “A. M. Turing Award”⁸ in the area of Computer Science or the “Bruce Gold Medal”⁹ in the area of Astronomy, are among the many examples of highly prestigious awards that are ignored in the Shanghai ranking.

Summarizing, the two criteria ALU and AWA are only very loosely connected with what they are trying to capture. Their evaluation furthermore involves arbitrary parameters and raises many difficult counting problems. Hence, these criteria are plagued by a significant imprecision and inaccurate determination (Bouyssou 1989; Roy 1988). Finally they are based on distinctions that do not cover every important scientific field.

Highly cited researchers

As stressed in van Raan (2005a), the most striking fact here is the complete reliance of the authors of the ranking on choices made by Thomson Scientific. Is the division of Science into 21 domains relevant? In view of Table 3, it is apparent that the choice of these 21 domains seems to favor medicine and biology. This may be a reasonable option for a commercial firm like Thomson Scientific, since these fields generate many “hot” papers.

⁸ Awarded every year since 1966 by the Association for Computing Machinery, see <http://www.awards.acm.org/homepage.cfm?awd=140>.

⁹ Awarded every year since 1898 by the Astronomical Society of the Pacific, see <http://www.phys-astro.sonoma.edu/bruceMedalists>.

Table 3 The 21 categories used by Thomson Scientific

21 categories used by Thomson Scientific	
Agricultural Sciences	Materials Science
Engineering	Plant & Animal Science
Neuroscience	Computer Science
Biology & Biochemistry	Mathematics
Geosciences	Psychology/Psychiatry
Pharmacology	Ecology/Environment
Chemistry	Microbiology
Immunology	Social Sciences, General
Physics	Economics & Business
Clinical Medicine	Molecular Biology & Genetics
Space Sciences	

Source: http://www.isihighlycited.com/isi_copy/Comm_newse04.htm

The fact that this choice is appropriate in order to evaluate universities would need to be justified and, unfortunately, the authors of the ranking remain silent on this point.

Moreover, these 21 categories do not have the same size. Speaking only in terms of the number of journals involved in each categories (but keep in mind that journals may have quite different sizes) they are indeed quite different. Space Science involves 57 journals, Immunology 120,..., Plant & Animal Science 887, Engineering 977, Social Science General 1299, and Clinical Medicine 1305 (source: http://www.isihighlycited.com/isi_copy/Comm_newse04.htm).

This criterion clearly uses Thomson Scientific citation counts. Bibliometricians have often stressed that these citation counts are somewhat imprecise. Indeed, the matching of cited papers involves “losses” (e.g., due to incorrect spelling or wrong page numbers). Van Raan (2005a) evaluates the average loss of citations to 7%, while it may be as high as 30% in certain fields. Bizarrely, Liu et al. (2005) answering these comments, simply did not acknowledge that criterion HiCi uses citation counts.

Finally, let us observe that Thomson Scientific uses a period of 20 years to determine the names of highly cited researchers in each category. Hence, in most categories, the persons in these lists are not particularly young and have often changed institutions several times during their careers.

Summarizing, combining the exclusive reliance on a division of Science into 21 categories suggested by Thomson Scientific, the use of a rather long period of reference, and the difficulties inherent to a precise counting of citations reveals that this criterion is only extremely loosely connected to the present ability of an institution to produce research with high impact.

Papers in *Nature* and *Science*

Probably the most surprising fact with this criterion is the weighting scheme for multiple authors (this is the usual rule in the “hard sciences”). With 100% for the corresponding author, 50% for the first author, 25% for the next author affiliation, and 10% for other author affiliations, one quickly sees that all papers published in *Nature* and *Science* do not have the same weight. A paper signed by many co-authors will have a greater weight than a paper signed by a single person (therefore it is in the interest of an institution that any paper published in *Nature* and *Science* is co-signed by many co-authors from the same

institution). We have to say that this seems highly counter-intuitive and even paradoxical. We should also mention that the problems of affiliation that we examine below are also present for this criterion.

Articles indexed by Thomson Scientific

As stressed in van Raan (2005a), the authors of the ranking entirely rely for the evaluation of this criterion on the Thomson Scientific databases. This raises a number of important problems.

First, the attribution of the papers to the right institution is far from being an easy task. The authors of the ranking solve it saying that “institutions or research organizations affiliated to a university are treated according to their own expression in the author affiliation of an article” (Liu and Cheng 2005, p. 134). This is likely to lead to many problems. For instance, it is well known that authors do not always pay much attention to the standardization of their affiliation when they publish a paper. The problem is especially serious when it comes to papers published by university hospitals (they often have a specific name that is distinct from the name of the university and have a distinct address, see van Raan 2005a; Vincke 2009). A similar phenomenon occurs when an institution has an official name that is not in English. A famous example is the difficulty to distinguish the *Université Libre de Bruxelles* from the *Vrije Universiteit Brussel*. Both are located in Brussels and have the same postal code. Both names are the same in English (Free University of Brussels). Hence this first problem is likely to cause much imprecision on the evaluation of criterion PUB. Attaching to each author a correct affiliation is a difficult task requiring a deep knowledge of the peculiarities of the institutional arrangements in each country.

Second, it is well known that the coverage of the Thomson Scientific database is in no way perfect (Adam 2002). The newly created SCOPUS database launched by Elsevier, has a vastly different coverage (although, clearly, the intersection between the two databases is not empty). Counting using Thomson Scientific instead of SCOPUS is a perfectly legitimate choice, provided that the impact of this choice on the results is carefully analyzed. This is not the case in the Shanghai ranking.

Third, it is also well known that the coverage of most citation database has a strong slant towards publications in English (see van Leeuwen et al. 2001; van Raan 2005a, for an analysis of the impact of this bias on the evaluation of German universities). Yet, there are disciplines (think of Law) in which publications in a language that is not the language of the country make very little sense. Moreover, there are whole parts of Science that do not use articles in peer-reviewed journals as the main media for the diffusion of research. In many parts of Social Science, books are still a central media, whereas in Engineering or Computer Science, conference proceedings dominate. The authors of the ranking have tried to correct for this bias against Social Sciences by multiplying by a factor 2 all papers indexed in the Social Science Citation Index. This surely goes in the right direction. But it is also quite clear that this coefficient is arbitrary and that the impact of varying it should be carefully analyzed.

Finally, we may also wonder why the authors of the ranking have chosen to count indexed papers instead of trying to measure the impact of the papers. Browsing through the Thomson Scientific databases quickly reveals that most of indexed papers are almost never cited and that a few of them concentrate most citations, this being true independently of the impact of the journal. The bibliometric literature has emphasized the importance of taking

the *impact* of research into account in order to produce relevant and meaningful indices (see, e.g., the works of Moed 2006; Moed et al. 1995; van Raan 1996, 2006).

Summarizing, criterion PUB raises several important problems and involves many arbitrary choices.

Productivity

Criterion Py consists in the “total score of the above five indicators divided by the number of Full Time Equivalent (FTE) academic staff”. It is ignored when this last number could not be obtained. Two main things have to be stressed here.

First this criterion is clearly affected by all the elements of imprecision and inaccurate determination analyzed above for the first five criteria. Moreover, the authors of the ranking do not fully detail which sources they use to collect information on the number of Full Time Equivalent (FTE) academic staff. The authors of the ranking rely here on a variety of sources (National Ministry of Education, National Bureau of Statistics, National Association of Universities and Colleges, National Rector’s Conference, see above). Since the notion of “member of academic staff” is not precisely defined and may be interpreted in several quite distinct ways (e.g., how to count invited or emeritus professors?), we have no reason to believe that information collected through these various sources is fully consistent and reliable.

Second, it is not 100% clear what is meant by the authors when they refer to the “total score of the above five indicators”. Are these scores first normalized? Are these scores weighted? (we suspect that this is the case). Using which weights? (we suspect that these weight are simply the weights of the first five indicators normalized to add up to 1).

A varying number of criteria

Institutions are evaluated in the Shanghai ranking using six criteria... but not all of them. In fact we have seen that there are several possible cases:

- institutions not specialized in Social Sciences and for which FTE academic staff data could be obtained are evaluated on 6 criteria: ALU, AWA, HiCi, N&S, PUB, and Py.
- institutions not specialized in Social Sciences and for which FTE academic staff data could not be obtained are evaluated on 5 criteria: ALU, AWA, HiCi, N&S, and PUB.
- institutions specialized in Social Sciences and for which FTE academic staff data could be obtained are evaluated on 5 criteria: ALU, AWA, HiCi, PUB, and Py.
- institutions specialized in Social Sciences and for which FTE academic staff data could not be obtained are evaluated on 4 criteria: ALU, AWA, HiCi, and PUB.

This raises many questions. First MCDM has rarely tackled the situation in which alternatives are not evaluated on the same family of criteria. This raises many interesting questions. For instance the right way to meaningfully “neutralize” a criterion does not seem to be entirely obvious. Second, the authors of the ranking do not make publicly available the list of institutions that they consider to be specialized in Social and Human Sciences. They neither give the precise list of institutions for which criterion Py could be computed. Hence not only the family of criteria varies but it is impossible to know which family is used to evaluate what.

A brief summary on criteria

We have seen that all criteria used by authors of the ranking are only loosely connected with what they intended to capture. The evaluation furthermore involves several arbitrary parameters and implies taking many micro-decisions that are not documented. In view of Fig. 1, we surely expect all these elements to quite severely impact the robustness of the results of the ranking. Quite unfortunately, since the authors of the ranking do not make “raw data” publicly available (a practice which does not seem to be fully in line with their announced academic motives), it is impossible to analyze the robustness of the final ranking with respect to these elements.

We have seen above that the authors claim that the ranking: uses “carefully selected objective criteria”, is “based on internationally comparable data that everyone can check”, and is such that “no subjective measures were taken”.

It seems now clear that the criteria have been chosen mainly based on availability, that each one of them is only loosely connected with what should be captured and that their evaluation involves the use of arbitrary parameters and arbitrary micro-decisions. The impact of these elements on the final result is not examined. The raw data that are used are not made publicly available so that they cannot be checked.

We would finally like to mention that there is a sizeable literature on the question of structuring objectives, associating criteria or attributes to objective, discussing the adequateness and consistency of a family of criteria. This literature has two main sources. The first one originates in the psychological literature (Ebel and Frisbie 1991; Ghiselli 1981; Green et al. 1988; Kerlinger and Lee 1999; Kline 2000; Nunally 1967; Popham 1981) has concentrated on the question of the *validity* and *reliability* and has permeated the bulk of empirical research in Social Sciences. The second originates from MCDM (Bouyssou 1990; Fryback and Keeney 1983; Keeney and Raiffa 1976; Keeney 1981, 1988a, b, 1992; Keeney and McDaniel 1999; Keeney et al. 1999; Roy 1996; Roy and Bouyssou 1993; von Winterfeldt and Edwards 1986) has concentrated on the question of the *structuration of objectives* and the question of the *construction of attributes or criteria* to measure the attainment of these objectives. It seems to have been mostly ignored by the authors of the ranking.¹⁰

¹⁰ Let us mention here several other problems with the criteria used by the authors of the ranking. First they have chosen to publish their ranking on an annual basis. This is probably a good choice if what is thought is media coverage. However, given the pace of most research programs, we cannot find any serious justification for such a periodicity. As observed in Gingras (2008), the ability of a university to produce excellent research, is not likely to change much from one year to another. Therefore, changes from one edition of the ranking to the next one are more likely to reflect random fluctuations than real changes. This is all the more true that several important points in the methodology and the criteria have changed over the years (Saisana and D’Hombres (2008), offer an overview of these changes). Second, the choice of an adequate period of reference to assess the “academic performance” of an institution is a difficult question. It has been implicitly answered by the authors of the ranking in a rather strange way. Lacking any clear analysis of the problem, they mix up in the model several very different time periods: one century for criteria ALU and AWA, 20 years for criterion HiCi, 5 years for criterion N&S, and 1 year for criterion PUB. There may be a rationale behind these choices but it is not made explicit by the authors of the ranking. As observed in van Raan (2006a, b), “academic performance” can mean two very different things: the prestige of an institution based on its past performances and its present capacity to attract excellent researchers. These two elements should not be confused. Third, five of the six criteria used by the authors of the ranking are counting criteria (prizes and medals, highly cited researchers, papers in N&S, papers indexed by Thomson Scientific). Hence, it should be no surprise that all these criteria are strongly linked to the size of the institution. As Zitt and Filliatreau (2006) have forcefully shown, using so many criteria linked to the size of the institution is the sign that *big is made beautiful*. Hence, the fact that criteria are highly correlated should not be a surprise. Although the authors of the ranking view this fact as a strong point of their approach, it is more likely to

An MCDM view on the Shanghai ranking

In the previous section, we have proposed a critical analysis of the criteria used by the authors of the ranking, mainly synthesizing the existing literature on the subject. We now turn to questions linked with the methodology used by the authors to aggregate these criteria. As far as we know these important aspects, that are well known in the literature on MCDM, have never been tackled so far in the literature related to “university rankings” (an early publication of some of these arguments was made in Vincke 2009, based on our joint work).

A rhetorical introduction

Suppose that you are giving a Master course on MCDM. The evaluation of students is based on an assignment consisting in proposing and justifying a particular MCDM technique on an applied problem. The subject given to your students this year consists in devising a technique that would allow to “rank order countries according to their ‘wealth’”. Consider now the case of three different students.

The first student has proposed a rather complex technique that has the following feature. The fact that country a is rank before or after country b does not only depend on the data collected on countries a and b but also with what happens with a third country c . Our guess is that you will find that the work of this student is of *very poor quality*. Indeed, the relative position of countries a and b should only depend upon their own performances. Although such a dependence on “irrelevant alternatives” may be rationalized in certain contexts (see Luce and Raiffa 1957; Sen 1993), we do not think that this is the case here.

Consider a second student that has proposed a simple technique that works as follows. For each country she has collected the GNP (Gross National Product) and the GNPpc (Gross National Product per capita) of this country. She then suggests to rank order the countries using a weighted average of the GNP and the GNPpc of each country. Our guess is that you will find that the work of this student is of *very poor quality*.

Either you want to measure the “total wealth” of a country and you should use the GNP or you want to measure the “average richness” of its inhabitants and you should use the GNPpc. Combining these two measures using a weighted average makes no sense: the first is a “production” measure, the second is a “productivity” measure. Taking α times production plus $(1 - \alpha)$ times productivity is something that you can compute but that has absolutely no meaning, unless, of course, if α is 0 or 1. The reader who is not fully convinced that this does not make sense is invited to test the idea using statistics on GNP and GNPpc that are widely available on the web: the results of such an experiment are quite perplexing.

Consider now a third student who has proposed a complex model but that has:

- not questioned the relevance of the task,
- not reflected on what “wealth” is and how it should be measured,
- not investigated the potential impacts of her work,

Footnote 10 continued

simply reflect the impact of size effects. Fourth, Since the criteria used by the authors of the ranking are linked with “academic excellence”, we should expect that they are poorly discriminatory between institutions that are not ranked among the top ones. A simple statistical analysis reveals that this is indeed the case, see Billaut et al. (2009).

- only used readily available information on the web without questioning its relevance and precision,
- has mixed this information with highly subjective parameters without investigating their influence of the results.

Clearly you will find that the work of this student is of *very poor quality*. Indeed, she has missed the entire difficulty of the subject reducing it to a mere number-crunching exercise.

We are sorry to say that the authors of the ranking do not seem to be in a much better position than any of our three students. We explain below why we think that they have, in their work, combined all what we have found to be highly questionable in the work of these three students.

The aggregation technique used is flawed

One of the first thing that is invariably taught in any basic course on MCDM is the following: if you aggregate several criteria using a weighted sum, the weights that are used should *not* be interpreted as reflecting the “importance” of the criteria. This may seem strange but is in fact very intuitive. Weights, or rather *scaling constants* as we call them in MCDM, are indeed linked to the *normalization* of the criteria. If normalization changes, weights should change. A simple example should help the reader not familiar with MCDM understand this point. Suppose that one of your criterion is a measure of length. You may choose to measure this criterion in meters, but you may also choose to measure it in kilometers. If you use the same weights in both cases, you will clearly end up with *absurd results*.

This has two main consequences. First, weights in a weighted sum cannot be assessed on the basis of a vague notion of “importance”. The comparison of weights used in a weighted sum do not reflect a comparison in terms of importance of the criteria. Indeed, if the weight of a criterion measured in meters is 0.3 this weight should be multiplied by 1 000 if you decide to measure it in kilometers. Therefore the comparison of this weight with the weights of other criteria does not reflect a comparison of importance (it may well happen that the weight of criterion length, when this criterion is measured in meters, is smaller than the weight of another criterion, while the opposite comparison will prevail when this criterion is measured in kilometers). This has many important consequences on the correct way to assess weights in a weighted sum (see Bouyssou et al. 2006; Keeney and Raiffa 1976). In any case, it does not make sense to ask someone directly for weights, in a weighted sum, based on a vague notion of “importance” (as the authors of the ranking do on their web site, see <http://www.arwu.org/rank/2004/Questionnaire.htm>). This also raises the problem on how the authors of the ranking have chosen their set of weights. They offer no clue on this point. It seems safe to consider that the weights have been chosen arbitrarily. The only rationale we can imagine for this choice is that, in the first version of the ranking, the authors used only five criteria with equal weights. Although the use of equal weights may be justified under certain circumstances (see Einhorn and Hogarth 1975), we have no reason to believe that they apply here.

The second and more devastating consequence is the following. If you change the normalization of the criteria, you should absolutely change the weights. If you do not do so, this amounts to changing the weights ... and you will end up with absurd results. Since, each year, the authors of the ranking normalize their criteria giving the score of 100 to the best scoring institution on each criterion, and, since each year the non-normalized score of the best scoring institution on this criterion is likely to change, the weights should change

Table 4 Weighted sum: example with equal weights

alternatives	g_1	g_2	g_1^n	g_2^n	Score	Rank
h	2,000	500	100.0	100.0	100.0	1
a	160	435	8.0	87.0	47.5	2
b	400	370	20.0	74.0	47.0	3
c	640	305	32.0	61.0	46.5	4
d	880	240	44.0	48.0	46.0	5
e	1,120	175	56.0	35.0	45.5	6
f	1,360	110	68.0	22.0	45.0	7
g	1,600	45	80.0	9.0	44.5	8

Table 5 Weighted sum with equal weights: h increases on g_2

alternatives	g_1	g_2	g_1^n	g_2^n	Score	Rank
h	2,000	700	100.00	100.00	100.00	1
a	160	435	8.00	62.14	35.07	8
b	400	370	20.00	52.86	36.43	7
c	640	305	32.00	43.57	37.79	6
d	880	240	44.00	34.29	39.14	5
e	1,120	175	56.00	25.00	40.50	4
f	1,360	110	68.00	15.71	41.86	3
g	1,600	45	80.00	6.43	43.21	2

each year so as to cope with this new normalization. But the authors of the ranking do not change the weights to reflect this change of normalization.¹¹

Let us illustrate what can happen with a simple example using two criteria. Let us consider the data in Table 4. In this table, eight alternatives (or institutions) a, b, c, d, e, f, g and h are evaluated on two criteria g_1 and g_2 (the average values that are used in this example roughly correspond to the average values for criteria PUB and $10 \times \text{HiCi}$). These criteria are normalized so as to give a score of 100 to the best scoring alternative on each criterion (here, h , as in Harvard, on both criteria). This defines the two normalized criteria g_1^n and g_2^n . For instance we have $g_2^n(e) = 35 = (175 \times 100)/500$. Let us aggregate these two criteria with a weighted sum using equal weights. This defines the ‘Score’ column in Table 4 (it is not necessary to normalize again the global score, since the score of h is already 100). If we use this global score to rank order the alternatives, we obtain the following ranking ($a \succ b$ means that a is preferred to b):

$$h \succ a \succ b \succ c \succ d \succ e \succ f \succ g.$$

Consider now a similar situation in which everything remains unchanged except that the performance of h on g_2 increases: it is now 700 instead of 500. This leads to the data in Table 5. The two criteria are again normalized so as to give a score of 100 to the best scoring alternative on each criterion (here again, h on both criteria). But because the score of h on g_2 has changed, this impacts *all* normalized scores on g_2^n . If you decide to aggregate the two normalized criteria using the same weights as before, you end up with the following ranking:

¹¹ Keeney (1992, p. 147) calls this the “most common critical mistake”.

$$h \succ g \succ f \succ e \succ d \succ c \succ b \succ a.$$

Observe that the modification of the score of h on g_2 has *inverted the ranking of all other alternatives*. The intuition behind this “paradox” should be clear. Since the score of h on g_2 has changed, we have changed the normalization of criterion g_2^n . Because the normalization has changed, the weight of this criterion should change if we want to be consistent: instead of using weights equal to 0.5 and to 0.5, we should now use different weights so as to reflect this change of normalization.

Observe that the failure to change weights when normalization changes has very strange effects besides the ones just mentioned. If an institution is weak on some criterion, so that a competitor is ranked just before it, its interest is that the best scoring alternative on this particular criterion improves its performance: if the weights are kept unchanged, this will mechanically decrease the importance of this criterion and will eventually allow it to be ranked before its competitor. Therefore if an institution is weak on some criterion, its interest is that the difference between its performance and the performance of the best scoring institution on this criterion increases!

Clearly, the above numerical examples have been chosen with care. They nevertheless show that it is impossible to assert that in the Shanghai ranking improving its performances on some criteria will necessarily lead to an improved position in the ranking. This should quite severely undermine the confidence we should have in the results of this ranking. A rather tortuous argument could be put forward in order to try to salvage the results of the Shanghai ranking saying that the data are such that, whatever the weights, the results are always the same. In view of Fig. 1, it seems clear that such an argument does not apply here.


Let us conclude with a final remark on the aggregation technique that is used. Even if the authors of the ranking had not fallen in the “normalization trap” explained above (and, clearly, there are very simple steps that could be taken to correct this point, e.g., choosing a normalization of the criteria that does not change every year), the weighted sum would remain a poor way to aggregate criteria. Almost all of Bouyssou et al. (2000) is devoted to examples explaining why this is so. Let us simply recall here the classical problem of the existence of *unsupported efficient alternatives*, in the parlance of MCDM. An alternative is said to be dominated if there is an alternative that has better evaluations on all criteria and a strictly better one on some criterion. An alternative is *efficient* if it is not dominated. It seems clear that all efficient alternatives are potentially interesting alternatives. A good aggregation technique should therefore allow any one of them to be ranked first with an adequate choice of parameters. Yet, with a weighted sum, there are efficient alternatives that cannot be ranked first, whatever the choice of weights.

Table 6 gives an example (taken from Bouyssou et al. 2000) of such a situation, using two criteria to be maximized. Observe that there are three efficient alternatives in this example: a , b , and c (alternative d is clearly dominated by all other alternatives). Intuitively, alternative c appears to be a very good candidate to be ranked first: it performs

Table 6 Weighted sum: unsupported efficient alternatives

	g_1	g_2
a	5	19
b	20	4
c	11	11
d	3	3

Fig. 2 Unsupported efficient alternatives



reasonably well on all criteria, while a (resp. b) is excellent on criterion 2 (resp. 1) but seems poor on criterion 1 (resp. 2). However, if the two criteria are aggregated using a weighted sum, it is impossible to find weights that would rank c on top. Indeed, suppose that they are weights α and $1 - \alpha$ that would allow to do so. Ranking c before a implies $11\alpha + 11(1 - \alpha) > 5\alpha + 19(1 - \alpha)$, i.e., $\alpha > 8/15 \approx 0.53$. Ranking c before b implies $11\alpha + 11(1 - \alpha) > 20\alpha + 4(1 - \alpha)$, i.e., $\alpha < 7/16 \approx 0.44$. Figure 2 shows that this impossibility is due to the fact that c is dominated by a convex combination of a and b . Recent research in MCDM have exhibited a wealth of aggregation techniques that do not have such a major deficiency (Belton and Stewart 2001; Bouyssou et al. 2000).

The aggregation technique that is used is nonsensical

Criteria ALU, AWA, HiCi, N&S, and PUB are *counting* criteria. It is therefore clear that they are globally linked to the ability of an institution to produce a large amount of good papers and good researchers. They capture, up to the remarks made in “[An analysis of the criteria](#)”, the “research potential” of an institution. This is semantically consistent. However, criterion Py is quite different. If the first five criteria capture “production” the last one captures “productivity”. But common sense and elementary economic analysis strongly suggest that taking a weighted average of production and productivity, although admissible from a purely arithmetic point of view, leads to a composite index that is meaningless (we use the word here in its ordinary sense and not in its measurement-theoretic sense, see Roberts 1979). The only argument that we can think of in favor of such a measure is that the weight of the last criterion is rather small (although, we have seen above that weights in a weighted sum should be interpreted with great care). Nevertheless, the very fact that production is mixed up with productivity seems to us highly problematic and indicates a poor reflection on the very meaning of what an adequate composite index should be. The projects of the authors of the ranking, as announced in Liu et al. (2005, p. 108), to build a ranking with a weight of 50% for the criterion Py, are a clear sign that this semantic problem has not been fully acknowledged by the authors of the ranking. This should severely undermines the confidence we can have in the results of the ranking.

Neglected structuring issues

When trying to build an evaluation model, good practice suggests (Bouyssou et al. 2000; JRC/OECD 2008) that the reflection should start with a number of simple but crucial questions:

1. What is the definition of the objects to be evaluated?
2. What is the purpose of the model? Who will use it?
3. How to structure objectives?
4. How to achieve a “consistent family of criteria”?
5. How to take uncertainty, imprecision, and inaccurate definition into account?

Concerning the last three questions, we have seen in “[An analysis of the criteria](#)” that the work of the authors of the ranking could be subjected to severe criticisms. This is especially true for the last question: since raw data are not made publicly available and the many micro-decisions that led to these data are not documented, it is virtually impossible to analyze the robustness of the proposed ranking. The partial analyses conducted in Saisana and D’Hombres (2008) show that this robustness is likely to be extremely weak.

Let us concentrate here on the first two questions keeping in mind a number of good practices for the construction of an evaluation model.

What is a “university”?

This question might sound silly to most of our readers coming from the US and the UK. However for a reader coming from continental Europe this question is not always an easy one. Let us take here the example of France, which is, admittedly, a particularly complex example. In France co-exist:

- Public universities (usually named *Universités*). What should be observed here is that the history of most of these universities has been long and somewhat erratic. After 1968, most of them were split into several smaller ones. Moreover, there are many newly created universities in France that are rather small and do not offer programs in all areas of Science and/or at all levels of the Bachelor–Master–Doctorate scale. Finally, when analyzing the French system, it should be kept in mind that these universities rarely attract the best students. They rather choose to enter the *Grandes Écoles* system. Tuition fees in these universities are generally small.
- *Grandes Écoles* (mainly in Engineering, Management and Political Science) are very particular institutions. They are usually quite small and most of them only grant Master degrees. They are highly selective institutions that are recruiting students after a nationwide competitive exam. They have a long tradition and a very active network of alumni. Only very few of them are actively involved in Doctoral programs. Tuition fees in *Grandes Écoles* vary a lot. Some of them are quite expensive (mostly management schools) while in some others, the fees are comparable to that of a public university. Finally, in some of them, (e.g., the *Écoles Normales Supérieures*), students are paid.
- Large public and private research institutes that may have common research centers, among them or with universities or *Grandes Écoles*. Among the public research centers we should mention: CNRS, INSERM (specialized in biomedical research), INRA (specialized in agricultural sciences) and INRIA (specialized in Computer Science). A very significant part of research in France is conducted in such institutes, although they have no student and grant no diploma. Moreover, there are large and renowned private research centers, the most famous one being the *Institut Pasteur* (many of the French Nobel prizes in Medicine are linked to it).

With such a complex institutional landscape, what should count as a university is far from being obvious. It seems that this was not obvious to the authors of the ranking, since they included in the first three editions (2003–2005) of the ranking the *Collège de France*,

an institution that has no student and grants no diploma: if such an institution can count as a university then almost all organizations can. The French situation is especially complex but is far from being exceptional. Germany and Italy have strong Public Research centers, besides their universities, too.

Any evaluation system should minimally start with a clear definition of the objects to be evaluated. Such a definition is altogether lacking in the Shanghai ranking.

What is a “good” university?

The authors of the ranking are interested in “world-class” universities. But, as they have not proposed a definition of what a university is, they do not offer a definition of what a “world class” university is. Nevertheless the criteria they use allow to implicitly define what they mean here. The only thing of importance is “excellence” in research. Moreover this excellence is captured using very particular criteria evaluated in a very particular way (see “[An analysis of the criteria](#)”). Why ignore research outputs such as patents, books or PhD theses? Why count papers instead of trying to measure impact?, etc.

Perhaps the most perplexing thing in the implicit definition of a world class university used by the authors of the ranking is that it mostly ignores *inputs* and *institutional constraints*.

Some universities have a more-or-less complete freedom to organize their governance, to hire and fire academic and non-academic staff, to decide on salaries, to select students, to decide on tuition fees. Some others have almost no freedom in all these respects (this is mostly the case for French universities). They cannot select students, they cannot decide on tuition fees, they are not fully involved in the selection of their academic staff, and firing someone is difficult. Given such differences in institutional constraints, should we simply ignore them, as is implicitly done in the ranking? This is only reasonable if one admits that there is “one best model” of a world-class university. This hypothesis would need detailed empirical justification that is not offered by the authors of the ranking.

Similarly the “inputs” consumed by institutions in their “scientific production process” are mostly ignored. The only input that is explicitly taken into account is the number of FTE academic staff, when it could be obtained. But there are many other important inputs that should be included, if one is to judge on the efficiency of a scientific production process. Let us simply mention here that tuition fees, funding (Harvard’s annual budget is over 3×10^9 USD in 2007, Harvard University 2007, p. 38; this is larger than the GDP of Laos), quality of campus, libraries (Harvard’s libraries possess over 15×10^6 volumes, Harvard University 2007), academic freedom to research and publish on any subject of interest, etc. are also very important ingredients in the success of a university. Ignoring all these inputs implies a shallow and narrow view on academic excellence.¹²

What is the purpose of the model? Who can profitably use it?

To us, the very interest of ranking “universities” is not obvious at all. Indeed, who can benefit from such a ranking?

Students and families looking for information are much more likely to be interested in a model that will evaluate *programs*. We are all aware of the fact that a good university may

¹² Let us remark that we disagree here with Principle 8 in International Ranking Expert Group (2006): a production process, whether it is or not scientific, cannot be analyzed without explicitly considering outputs and inputs.

be especially strong in some areas and quite weak in others. Moreover, it seems clear that (although we realize that each family might want to consider its child as a future potential Nobel Prize winner) students and families are likely to be interested in rather trivial things such as tuition fees, quality of housing, sports facilities, quality of teaching, reputation of the program in firms, average salaries after graduation, strength of the alumni network, campus life, etc. For an interesting system offering such details, we refer to Berghoff and Federkeil (2009) and Centre for Higher Education Development (2008).

Recruiters are likely to be little impressed by a few Nobel prizes granted long ago to members of a given department, if they consider recruiting someone with a Master degree coming from a totally different department. Clearly, they will be mostly interested in the “employability” of students with a given degree. Besides the criteria mentioned above for students and families, things like the mastering of foreign languages, international experience, internships, etc. are likely to be of central importance to them.

Likewise, a global ranking of universities is quite unlikely to be of much use to deans and rectors willing to work towards an increase in quality. Clearly, managers of a university will be primarily interested in the identification of weak and strong departments, the identification of the main competitors, and the indication of possible directions for improvement. Unless they have a contract explicitly specifying that they have to increase the position of their institution in the Shanghai ranking (as astonishing as it may sound, this has happened), we do not see how a ranking of an institution *as a whole* can lead to a useful management tool.

Finally, political decision makers should be primarily interested in an evaluation system that would help them decide on the efficiency of the *higher education system of a country*. If a country has many good medium-sized institutions, it is unlikely that many of them will be standing high in the Shanghai ranking. But this does not mean that the system as a whole is inefficient. Asking for “large” and “visible” institutions in each country may involve quite an inefficient use of resources. Unless the authors of the ranking can produce clear empirical evidence that scientific potential is linked with size and that medium-sized institutions simply cannot produce valuable research, we do not understand why all this may interest political decision makers, except, of course, to support other strategic objectives.

Good evaluation practices

As detailed in Bouyssou et al. (2000, 2006), there are a number of good practices that should be followed when building an evaluation model. We want only to mention two of them here.

The first one is fairly obvious. If you evaluate a person or an organization, you should allow that person or organization to check the data that are collected on her/it. This seems quite obvious. Not doing so, inevitably leads to a bureaucratic nightmare in which each one is evaluated based on data that remain “behind the curtain”. We have seen that this elementary good practice has been forgotten by the authors of the ranking: since raw data are not made publicly available, it is impossible for the institutions that are evaluated to check them.

The second good practice we would like to mention is less clear-cut, but is nevertheless crucial. When an evaluation system is conceived, its creators should not expect the persons or the organizations that are evaluated to react passively to the system. This is the baseline of any introductory management course. Persons and organizations will adapt their behavior, consciously or not, in reaction to the evaluation system. This feedback is

inevitable and perverse effects due to such adaptations are inescapable (all this has been well documented in the management literature, Berry 1983; Boudon 1979; Dörner 1996; Hatchuel and Molet 1986; Mintzberg 1979; Moisdon 2005; Morel 2002). A good practice is therefore the following. Try to anticipate the most obvious perverse effects that can be generated by your evaluation system. Try to conceive a system in which the impacts of the most undesirable perverse effects are reduced. It does not seem that the authors of the ranking have followed this quite wise advice. The only words of wisdom are here that “Any ranking exercise is controversial, and no ranking is absolutely objective” and that “People should be cautious about any ranking and should not rely on any ranking either, including the ‘Academic Ranking of World Universities’”. Instead, people should use rankings simply as one kind of reference and read the ranking methodology carefully before looking at the ranking lists” (ARWU 2003–2009). Sure enough. But beyond that, we surely expect the developers of an evaluation system to clearly analyze the potential limitations of what they have created in order to limit, as far as possible, its illegitimate uses and the authors of the ranking remain silent on this point.

Suppose that you manage a university and that you want to increase your position in the ranking. This is simple enough. There are vast areas in your university that do not contribute to your position in the ranking. We can think here of Law, Humanities and most Social Sciences. Drop all these fields. You will surely save much money. Use this money to buy up research groups that will contribute to your position in the ranking. Several indices provided by Thomson Scientific are quite useful for this purpose: after all, the list of the potential next five Nobel prizes in Medicine is not that long. And, anyway, if the group is not awarded the prize, it will publish much in journals that count in the ranking and its members are quite likely to be listed among the highly cited researchers in the field. This tends to promote a view of Science that much resembles professional sports in which a few wealthy teams compete worldwide to attract the best players. We are not fully convinced that this is the best way to increase human knowledge, to say the least.

Manipulations are almost as simple and as potentially damaging for governments. Let us take for example the case of the French government, since we have briefly evoked above the complex organization of the French higher education system. Most French universities were split in several smaller parts in the early seventies. The idea was then to create organizations that would be easier to manage. Indeed, the venerable *Université de Paris* gave rise to no less than 13 new universities. But we have seen that this is surely detrimental in the ranking. So you should give these universities strong incentives to merge again. Neglecting the impact of the last criterion, a simple calculation shows that merging the universities in Paris that are mainly oriented towards “hard sciences” and Medicine (there is clearly no interest to merge with people doing such futile things as Law, Social Sciences and Humanities), i.e., Paris 5, 6, 7 and 11 (these are not the official names but their most common names), would lead (using the data from the 2007 Shanghai ranking) to an institution that would roughly be at the level of Harvard University. Bingo! You are not spending one more Euro, you have surely not increased the scientific production and potential of your country, you have created a huge organization that will surely be rather difficult to manage ... but you have impressively increased the position of France in the Shanghai ranking. Can you do even more? Sure, you can. Public research centers, although quite efficient, count for nothing in the ranking. You can surely suppress them and transfer all the money and persons to the huge organization you have just created. Then, you will surely end up much higher than Harvard University... No need to say that all these manipulations may lead, in the long term, to disastrous results.

Where do we go from here?

Let us now summarize our observations on the Shanghai ranking and try to draw some conclusions based on our findings, both on a scientific and a more strategic level.

An assessment of the Shanghai ranking

In what was probably the first serious analysis of the Shanghai ranking, van Raan (2005a, p. 140) stated that “From the above considerations we conclude that the Shanghai ranking should not be used for evaluation purposes, even not for benchmarking” and that “The most serious problem of these rankings is that they are considered as ‘quasi-evaluations’ of the universities considered. This is absolutely unacceptable”. We surely agree. The rather radical conclusions of van Raan were mainly based on bibliometric considerations, to which the authors of the ranking proved unable to convincingly answer (Liu et al. 2005; van Raan 2005b).

Our own analysis adopted a point of view that reflects our slant towards MCDM. Adding an MCDM point of view to the bibliometric analysis of van Raan (2005a) inevitably leads to an even more radical conclusion. Indeed, we have seen all criteria used by authors of the ranking are only loosely connected with what they intended to capture. The evaluation of these criteria involves several arbitrary parameters and many micro-decisions that are not documented. Moreover, we have seen that the aggregation method that is used is flawed and nonsensical. Finally, the authors of the ranking have paid almost no attention to fundamental structuring issues. Therefore, it does not seem unfair to say that *the Shanghai ranking is a poorly conceived quick and dirty exercise*. Again any of our MCDM student that would have proposed such a methodology in her Master’s Thesis would have surely failed according to our own standards.

What can be done?

An optimistic point of view would be that, after having read our paper, the authors of the ranking would decide to immediately stop their work, apologizing for having created so much confusion in the academic world, and that all political decision makers would immediately stop using “well known international rankings” as means to promote their own strategic objectives. However, we live in the real world and our bet is that this will not happen. Since the authors of Shanghai ranking more or less decided to ignore the point of view of van Raan (2005a) (as well as the ones expressed in further critical papers mentioned in “Introduction”), it is much likely that they will ignore ours. Therefore, we expect that they will continue for a long time to produce an annual ranking. Also, we should not expect too much of the willingness of political decision makers to abandon easy-to-use arguments that look striking enough in the general media. Therefore, we will have to live in a world in which extremely poor rankings are regularly published and used. What can be done then? Well, several things.

The first, and the more easy one, should be to stop being naive. “What is the best car in the world?”, “Where is the most pleasant city in Europe?”, “What is the best wine in the world?”, etc. All these questions may be interesting if your objective is to sell many copies of a newspaper or a book. However, it is clear that all these questions are meaningless unless they are preceded by a long and difficult structuring work. Clearly the “best car in the world” is a meaningless concept unless you have identified stakeholders, structured their objectives, studied the various ways in which attributes can be conceived to measure

the attainment of these objectives, applied meaningful procedures to aggregate this information and performed an extensive robustness analysis. Doing so, you might arrive at a model that can really help someone choose a car, or, alternatively, help a government to prepare new standards for greenhouse gas emissions. Without this work, the question is meaningless. If the question is meaningless for cars, should we expect a miracle when we turn to incredibly more complex objects such as universities? Certainly not. There is no such thing as a “best university” *in abstracto*. Hence, a first immediate step that we suggest is the following. Stop talking about these “all purpose rankings”. They are meaningless. Lobby in our own institution so that these rankings are never mentioned in institutional communication. This is, of course, especially important for our readers “lucky” enough to belong to institutions that are well ranked. They should resist the temptation of saying or thinking “there is almost surely something there” and stop using the free publicity offered by these rankings.

Since the production of poor rankings is unlikely to stop, a more proactive way to fight them is to produce many alternative rankings that produce vastly different results. It is not of vital importance that these new rankings are much “better” in some sense than the Shanghai ranking. Their main usefulness will be to “dilute” its devastating effects. A very interesting step in the direction was taken in (ENSMP 2007). The *École Nationale Supérieure des Mines de Paris* (ENSMP) is a French *Grande École*, being very prestigious in France. Its size is such that it is clear that it will never appear in good position in the Shanghai ranking. Hence, the ENSMP has decided to produce an alternative ranking. It can be very simply explained since it is based on a single criterion: the number of alumni of an institution having become the CEO of one of the top 500 leading companies as identified by the *Fortune* magazine (we refer to ENSMP 2007, for details on how this number is computed). We do not regard this ranking as very attractive. Indeed, the performances of the various institutions are based on things that happened long ago (it is quite unusual to see someone young becoming the CEO of a very large company). Therefore this number has little relation with the present performance of the institution. Moreover, this criterion is vitally dependent upon the industrial structures of the various countries in the world (institutions coming from countries in which industry is highly concentrated have a clear advantage) and “network effects” will have a major impact on it (we know that these effects are of utmost importance to understand the French *Grandes Écoles* system). Yet, we do not consider that the ENSMP ranking is much worse than the Shanghai ranking. On the contrary, it quite clearly points out the arbitrariness of many of the criteria used by the authors of the ranking. Finally and quite interestingly, the ENSMP ranking gives results that are *vastly different* from the Shanghai ranking. The top 10 institutions in this ranking have no less than 5 French institutions (and, among these 5, 3 are not even present in the top 500 of the Shanghai ranking). Diluting the effects of the Shanghai ranking clearly calls for *many more rankings* of this kind.

Why don’t you propose your own ranking?

We have been rather critical on the Shanghai ranking. The reader may legitimately wonder why we do not now suggest another, better, ranking. We will refrain here from doing so for various reasons.

First, this is not our main field of research. We are not experts in education systems. We are not experts in bibliometry. Contrary to the authors of the ranking, we do think that a decent ranking should obviously be the product of a joint research team involving experts in evaluation, education systems and bibliometry.

Second, we do not underestimate the work involved. Although we think that the Shanghai ranking is of poor quality, we realize that producing it is a *huge* task. A better ranking will inevitably involve even more work, based on the reasonable assumption that there is no free lunch.

Third, we are not at all convinced that this exercise is really useful. We have explained above why we would tend to rank programs or national education systems instead of universities. We have also given several clues on what should be done and what should absolutely not be done. We hope that these guidelines will be helpful for readers willing to take up to this task.

It seems likely that the combination of Operational Research (OR) with sophisticated bibliometric techniques will offer good promises for meaningful evaluation systems. Because the issue of bibliometric techniques has been extensively dealt with elsewhere (Moed 2006; Moed et al. 1995; van Raan 2005c, 2006; Zitt et al. 2005), let us mention here a few potentially useful OR techniques.

We have already mentioned how the absence of a minimal knowledge of aggregation techniques and their properties, as studied in MCDM, may vitiate an evaluation technique. We would like here to add that OR has also developed quite sophisticated tools to help structuring problems (Checkland 1981; Checkland and Scholes 1990; Eden 1988; Eden et al. 1983; Friend and Hickling 1987; Ostanello 1990; Rosenhead 1989) and to use this work together with sophisticated aggregation tools (Ackermann and Belton 2006; Bana e Costa et al. 1999; Belton et al. 1997; Montibeller et al. 2008; Phillips and Bana e Costa 2007). A parallel development concerns methods designed to assess the efficiency of “decision making units” that transform several inputs into several outputs, known as Data Envelopment Analysis (DEA) (Banker et al. 1984; Charnes et al. 1978; Cherchye et al. 2008; Cook and Zhu 2008; Cooper et al. 1999; Norman and Stoker 1991). The potential of such techniques for building evaluation models in the area of higher education has been already recognized (Bougnol and Dulá 2006; Johnes 2006; Leitner et al. 2007; Turner 2005, 2008). We suspect that a wise combination of the above two approaches is quite likely to lead to interesting evaluation models. It could well be combined with the promising interactive approach developed by the Center for Higher Education in Germany (Berghoff and Federkeil 2009; Centre for Higher Education Development 2008) in which users have the freedom to select the criteria they are willing to consider.¹³

Acknowledgements We wish to thank Florence Audier, Ghislaine Filliatreau, Thierry Marchant, Michel Zitt, and an anonymous referee for their useful comments on an earlier draft of this text.

References

- Ackermann, F., & Belton, V. (2006). Problem structuring without workshops? Experiments with distributed interaction in a PSM. *Journal of the Operational Research Society*, 58, 547–556.
- Adam, D. (2002). Citation analysis: The counting house. *Nature*, 415(6873), 726–729.
- Academic Ranking of World Universities (ARWU). (2003–2009). Shanghai Jiao Tong University, Institute of Higher Education, <http://www.arwu.org>.
- Bana e Costa, C. A., Ensslin, L., Corrêa, É. C., & Vansnick, J.-C. (1999). Decision support systems in action: Integrated application in a multicriteria decision aid process. *European Journal of Operational Research*, 113, 315–335.

¹³ During the preparation of this text, a European research consortium (CHERPA) won a call for tenders launched by the European Union on the question of university rankings. We wish much success to this international consortium and we, of course, hope that they will find parts of this text useful to them.

- Banker, R. D., Charnes, A., & Cooper, W. W. (1984). Some models for estimating technical and scale inefficiencies in data envelopment analysis. *Management Science*, 30(9), 1078–1092.
- Belton, V., & Stewart, T. J. (2001). *Multiple criteria decision analysis: An integrated approach*. Dordrecht: Kluwer.
- Belton, V., Ackermann, F., & Shepherd, I. (1997). Integrated support from problem structuring through alternative evaluation using COPE and V•I•S•A. *Journal of Multi-Criteria Decision Analysis*, 6, 115–130.
- Berghoff, S., & Federkeil, G. (2009). The CHE approach. In: D. Jacobs & C. Vermandele (Eds.), *Ranking universities* (pp. 41–63). Brussels: Édition de l'Université de Bruxelles.
- Berry, M. (1983). *Une technologie invisible ? Le rôle des instruments de gestion dans l'évolution des systèmes humains*. Mémoire Centre de Recherche en Gestion. École Polytechnique.
- Billaut, J.-C., Bouyssou, D., & Vincke, Ph. (2009). *Should you believe in the Shanghai ranking? An MCDM view*. Cahier du LAMSADE # 283, LAMSADE, <http://www.hal.archives-ouvertes.fr/hal-00388319/en/>.
- Boudon, R. (1979). *Effets pervers et ordre social*. Paris: PUF.
- Bougnol, M.-L., & Dulá, J. H. (2006). Validating DEA as a ranking tool: An application of DEA to assess performance in higher education. *Annals of Operations Research*, 145, 339–365.
- Bourdin, J. (2008). *Le défi des classements dans l'enseignement supérieur*. Rapport au Sénat 442, République française.
- Bouyssou, D. (1989). Modelling inaccurate determination, uncertainty, imprecision using multiple criteria. In: A. G. Lockett & G. Islei (Eds.), *Improving decision making in organisations* (LNEMS 335, pp. 78–87). Berlin: Springer-Verlag.
- Bouyssou, D. (1990). Building criteria: A prerequisite for MCDA. In: C. A. Bana e Costa (Ed.), *Readings in multiple criteria decision aid* (pp. 58–80). Springer-Verlag, Heidelberg.
- Bouyssou, D., Marchant, Th., Pirlot, M., Perny, P., Tsoukiàs, A., & Vincke, Ph. (2000). *Evaluation and decision models: A critical perspective*. Dordrecht: Kluwer.
- Bouyssou, D., Marchant, Th., Pirlot, M., Tsoukiàs, A., & Vincke, P. (2006). *Evaluation and decision models: Stepping stones for the analyst*. New York: Springer.
- Brooks, R. L. (2005). Measuring university quality. *The Review of Higher Education*, 29(1), 1–21.
- Buela-Casal, G., Gutiérrez-Martínez, O., Bermúdez-Sánchez, M. P., & Vadillo-Muñoz, O. (2007). Comparative study of international academic rankings of universities. *Scientometrics*, 71(3), 349–365.
- CHE. (2008). Centre for higher education development ranking. Technical report, <http://www.che.de>.
- Charnes, A., Cooper, W. W., & Rhodes, E. (1978). Measuring the efficiency of decision making units. *European Journal of Operational Research*, 2, 429–444. Correction: *European Journal of Operational Research*, 3:339.
- Checkland, P. (1981). *Systems thinking, systems practice*. New York: Wiley.
- Checkland, P., & Scholes, J. (1990). *Soft systems methodology in action*. New York: Wiley.
- Cherchye, L., Moesen, W., Rogge, N., van Puyenbroeck, T., Saisana, M., Saltelli, A., et al. (2008). Creating composite indicators with DEA and robustness analysis: The case of the technology achievement index. *Journal of Operational Research Society*, 59, 239–251.
- Cook, W. A., & Zhu, J. (2008). *Data envelopment analysis: Modeling operational processes and measuring productivity*. CreateSpace.
- Cooper, W. W., Seiford, L. M., & Tone, K. (1999). *Data envelopment analysis. A comprehensive text with models, applications, references and DEA-solver software*. Boston: Kluwer.
- Dalsheimer, N., & Despréaux, D. (2008). Analyses des classements internationaux des établissements d'enseignement supérieur. *Éducation & formations*, 78, 151–173.
- Desbois, D. (2007). Classement de Shanghai: peut-on mesurer l'excellence académique au niveau mondial? *La revue trimestrielle du réseau Écristin*, 67, 20–26.
- Dill, D., & Soo, M. (2005). Academic quality, league tables, and public policy: A cross-national analysis of university ranking systems. *Higher Education*, 49, 495–533.
- Dörner, D. (1996). *The logic of failure*. Jackson: Perseus Books.
- Ebel, R. L., & Frisbie, D. A. (1991). *Essentials of educational measurement*. New York: Prentice-Hall.
- Eden, C. (1988). Cognitive mapping. *European Journal of Operational Research*, 36, 1–13.
- Eden, C., Jones, S., & Sims, D. (1983). *Messing about in problems*. Oxford: Pergamon Press.
- Einhorn, H. J., & Hogarth, R. (1975). Unit weighting schemes for decision making. *Organizational Behavior and Human Performance*, 13, 171–192.
- Enserink, M. (2007). Who ranks the university rankers? *Science*, 317(5841), 1026–1028.
- ENSMP. (2007). *Professional ranking of world universities*. Technical report, École Nationale Supérieure des Mines de Paris (ENSMSP).
- Florian, R. (2007). Irreproducibility of the results of the Shanghai academic ranking of world universities. *Scientometrics*, 72, 25–32.

- Friend, J. K., & Hickling, A. (1987). *Planning under pressure: The strategic choice approach*. New York: Pergamon Press.
- Fryback, D. G., & Keeney, R. L. (1983). Constructing a complex judgmental model: An index of trauma severity. *Management Science*, 29, 869–883.
- Ghiselli, E. E. (1981). *Measurement theory for the behavioral sciences*. San Francisco: W. H. Freeman.
- Gingras, Y. (2008). Du mauvais usage de faux indicateurs. *Revue d'Histoire Moderne et Contemporaine*, 55(4bis), 67–79.
- Green, P. E., Tull, D. S., & Albaum, G. (1988). *Research for marketing decisions*. NJ: Englewood Cliffs. Harvard University.
- Harvard University. (2007). *Harvard university fact book, 2006–2007*. Technical report, Harvard University.
- Hatchuel, A., & Molet, H. (1986). Rational modelling in understanding and aiding human decision making: About two case studies. *European Journal of Operational Research*, 24, 178–186.
- CHERI/HEFCE. (2008). *Counting what is measured or measuring what counts? league tables and their impact on higher education institutions in England*. Report to HEFCE by the Centre for Higher Education Research and Information (CHERI) 2008/14, Open University, and Hobsons Research.
- International Ranking Expert Group. Berlin principles on ranking of higher education institutions. Technical report, CEPES-UNESCO.
- Ioannidis, J. P. A., Patsopoulos, N. A., Kavvoura, F. K., Tatsioni, A., Evangelou, E., Kouri, I., et al. (2007). International ranking systems for universities and institutions: A critical appraisal. *BioMed Central*, 5(30).
- Johnes, J. (2006). Measuring efficiency: A comparison of multilevel modelling and data envelopment analysis in the context of higher education. *Bulletin of Economic Research*, 58(2), 75–104.
- JRC/OECD. (2008). *Handbook on constructing composite indicators. methodology and user guide*. Technical report, JRC/OECD, OECD Publishing. ISBN 978-92-64-04345-9.
- Kävelmark, T. (2007). *University ranking systems: A critique*. Technical report, Irish Universities Quality Board.
- Keeney, R. L. (1981). Measurement scales for quantifying attributes. *Behavioral Science*, 26, 29–36.
- Keeney, R. L. (1988a). Structuring objectives for problems of public interest. *Operations Research*, 36, 396–405.
- Keeney, R. L. (1988b). Building models of values. *European Journal of Operational Research*, 37(2), 149–157.
- Keeney, R. L. (1992). *Value-focused thinking. A path to creative decision making*. Cambridge: Harvard University Press.
- Keeney, R. L., & McDaniel, T. L. (1999). Identifying and structuring values to guide integrated resource planning at BC gas. *Operations Research*, 47(5), 651–662. September–October.
- Keeney, R. L., & Raiffa, H. (1976). *Decisions with multiple objectives: Preferences and value tradeoffs*. New York: Wiley.
- Keeney, R. L., Hammond, J. S., & Raiffa, H. (1999). *Smart choices: A guide to making better decisions*. Boston: Harvard University Press.
- Kerlinger, F. N., & Lee, H. B. (1999). *Foundations of behavioral research* (4th ed.). New York: Wadsworth Publishing.
- Kivinen, O., & Hedman, J. (2008). World-wide university rankings: A scandinavian approach. *Scientometrics*, 74(3), 391–408.
- Kline, P. (2000). *Handbook of psychological testing* (2nd edn). New York: Routledge.
- Leimer, K.-H., Prikoszovits, J., Schaffhauser-Linzatti, M., Stowasser, R., & Wagner, K. (2007). The impact of size and specialisation on universities' department performance: A DEA analysis applied to Austrian universities. *Higher Education*, 53(4), 517–538.
- Liu, N. C. (2009). The story of academic ranking of world universities. *International Higher Education*, 54, 2–3.
- Liu, N. C., & Cheng, Y. (2005). The academic ranking of world universities. *Higher Education in Europe*, 30(2), 127–136.
- Liu, N. C., Cheng, Y., & Liu, L. (2005). Academic ranking of world universities using scientometrics: A comment to the "fatal attraction". *Scientometrics*, 64(1), 101–109.
- Luce, R. D., & Raiffa, H. (1957). *Games and decisions*. New York: Wiley.
- Marginson, S. (2007). *Global university rankings: Where to from here?* Technical report, Asia-Pacific Association for International Education.
- Mintzberg, H. (1979). *The structuring of organizations*. Englewood Cliffs: Prentice Hall.
- Moed, H. F., De Bruin, R. E., & van Leeuwen, T. N. (1995). New bibliometric tools for the assessment of national research performance: Database description, overview of indicators and first applications. *Scientometrics*, 33, 381–422.

- Moed, H. M. (2006). *Bibliometric rankings of world universities*. Technical Report 2006-01, CWTS, Leiden University.
- Moisdon, J.-C. (2005). Vers des modélisations apprenantes? *Économies et Sociétés. Sciences de Gestion*, 7, 569–582.
- Montibeller, G., Ackermann, F., Belton, V., & Ensslin, L. (2008). Reasoning maps for decision aiding: An integrated approach for problem structuring and multi-criteria evaluation. *Journal of the Operational Research Society*, 59, 575–589.
- Morel, C. (2002). *Les Décisions Absurdes*. Paris: Bibliothèque des Sciences Humaines. Gallimard.
- Norman, M., & Stoker, B. (1991). *Data envelopment analysis: The assessment of performance*. London: Wiley.
- Nunally, J. C. (1967). *Psychometric theory*. New York: McGraw-Hill.
- Ostanello, A. (1990). Action evaluation and action structuring: Different decision aid situations reviewed through two actual cases. In C. A. Bana e Costa (Ed.), *Readings in multiple criteria decision aid* (pp. 36–57). Springer-Verlag, Berlin.
- Phillips, L. D., & Bana e Costa, C. A. (2007). Transparent prioritisation, budgeting and resource allocation with multi-criteria decision analysis and decision conferencing. *Annals of Operations Research*, 154, 51–68.
- Popham, W. J. (1981). *Modern educational measurement*. New York: Prentice-Hall.
- Roberts, F. S. (1979). *Measurement theory with applications to decision making, utility and the social sciences*. Reading: Addison-Wesley.
- Rosenhead, M. J. (1989). *Rational analysis for a problematic world*. New York: Wiley.
- Roy, B. (1988). Main sources of inaccurate determination, uncertainty and imprecision in decision models. In: B. Munier & M. Shakun (Eds.), *Compromise, Negotiation and group decision* (pp. 43–67). Dordrecht: Reidel.
- Roy, B. (1996). *Multicriteria methodology for decision aiding*. Kluwer, Dordrecht. Original version in French “*Méthodologie multicritère d’aide à la décision*”, Economica, Paris, 1985.
- Roy, B., & Bouyssou, D. (1993). *Aide multicritère à la décision : méthodes et cas*. Paris: Economica.
- Saisana, M., & D’Hombres B. (2008). *Higher education rankings: Robustness issues and critical assessment. How much confidence can we have in higher education rankings?* Technical Report EUR 23487 EN 2008, IPSC, CRELL, Joint Research Centre, European Commission.
- Sen, A. K. (1993). Internal consistency of choice. *Econometrica*, 61, 495–521.
- Stella, A., & Woodhouse, D. (2006). *Ranking of higher education Institutions*. Technical report, Australian Universities Quality Agency.
- THES. (2008). Times higher education ranking supplement.
- T’kindt, V., & Billaut, J.-C. (2006). *Multicriteria scheduling* (2nd revised edition). Berlin: Springer Verlag.
- Turner, D. (2005). Benchmarking in universities: League tables revisited. *Oxford Review of Education*, 31(3), 353–371.
- Turner, D. (2008). World university rankings. *International Perspectives on Education and Society*, 9, 27–61.
- van Leeuwen, T. N., Moed, H. F., Tijssen, R. J. W., Visser, M. S., & van Rann, A. F. J. (2001). Language biases in the coverage of the *Science Citation Index* and its consequences for international comparisons of national research performance. *Scientometrics*, 51(1), 335–346.
- van Raan, A. F. J. (1996). Advanced bibliometric methods as quantitative core of peer review based evaluation and foresight exercises. *Scientometrics*, 36, 397–420.
- van Raan, A. F. J. (2005a). Fatal attraction: Ranking of universities by bibliometric methods. *Scientometrics*, 62, 133–145.
- van Raan, A. F. J. (2005b). Reply to the comments of Liu et al. *Scientometrics*, 64(1), 111–112.
- van Raan, A. F. J. (2005c). Measurement of central aspects of scientific research: Performance, interdisciplinarity, structure. *Measurement: Interdisciplinary Research and Perspectives*, 3(1), 1–19.
- van Raan, A. F. J. (2006). Challenges in the ranking of universities. In: J. Sadlak & N. C. Liu (Eds.), *World-class university and ranking: Aiming beyond status* (pp. 81–123). Bucharest: UNESCO-CEPES. ISBN 92-9069-184-0.
- Vincke, Ph. (2009). University rankings. In: D. Jacobs & C. Vermandele (Eds.), *Ranking universities*, (pp. 11–26) Brussels: Édition de l’Université de Bruxelles.
- von Winterfeldt, D., & Edwards, W. (1986). *Decision analysis and behavioral research*. Cambridge: Cambridge University Press.
- Zitt, M., & Filliatreau, G. (2006). Big is (made) beautiful: Some comments about the Shanghai-ranking of world-class universities. In: J. Sadlak & N. C. Liu (Eds.), *World-class university and ranking: Aiming beyond status* (pp. 141–160). Bucharest: UNESCO-CEPES (ISBN 92-9069-184-0).
- Zitt, M., Ramanana-Rahary, S. & Bassecoullard, E. (2005). Relativity of citation performance and excellence measures: From cross-field to cross-scale effects of field-normalisation. *Scientometrics*, 63(2), 373–401.